

Introduction to Machine Learning

Day I


Labatt Impact Lab Bootcamp

Alex Olson

- Studied Artificial Intelligence at the University of Edinburgh
- Came to UofT in 2018 to research applications of machine learning
- Published research applying machine learning to a wide range of disciplines: urban studies, cardiovascular surgery, construction...
- Helped to develop the very first ML Bootcamp in 2019
- Now: research associate at CARTE

Teaching Assistants

- Christina Seo and Jesse Ward-Bond
- Graduate students under Timothy Chan
- Researching Applied Optimization and Machine Learning




Artificial Intelligence

Getting computers to behave intelligently:

- Perform **non-trivial tasks** as well as humans do
- Perform **tasks that even humans struggle with**

Many sub-goals:

- Perception
- Reasoning
- Control
- Planning



My poker face: AI wins multiplayer game for first time

Pluribus wins 12-day session of Texas hold'em against some of the world's best human players



Speech Recognition

Perception + Reasoning



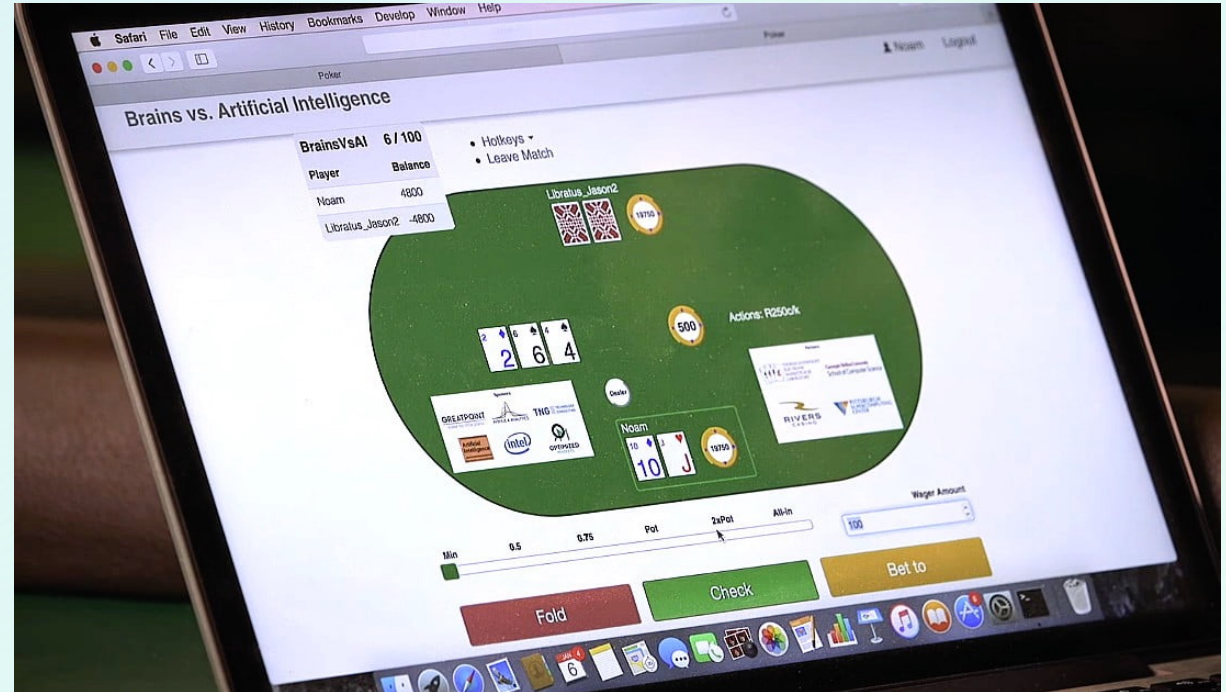

Autonomous Driving

Perception + **Reasoning**
Control + **Planning**



Game Playing

Reasoning + Planning

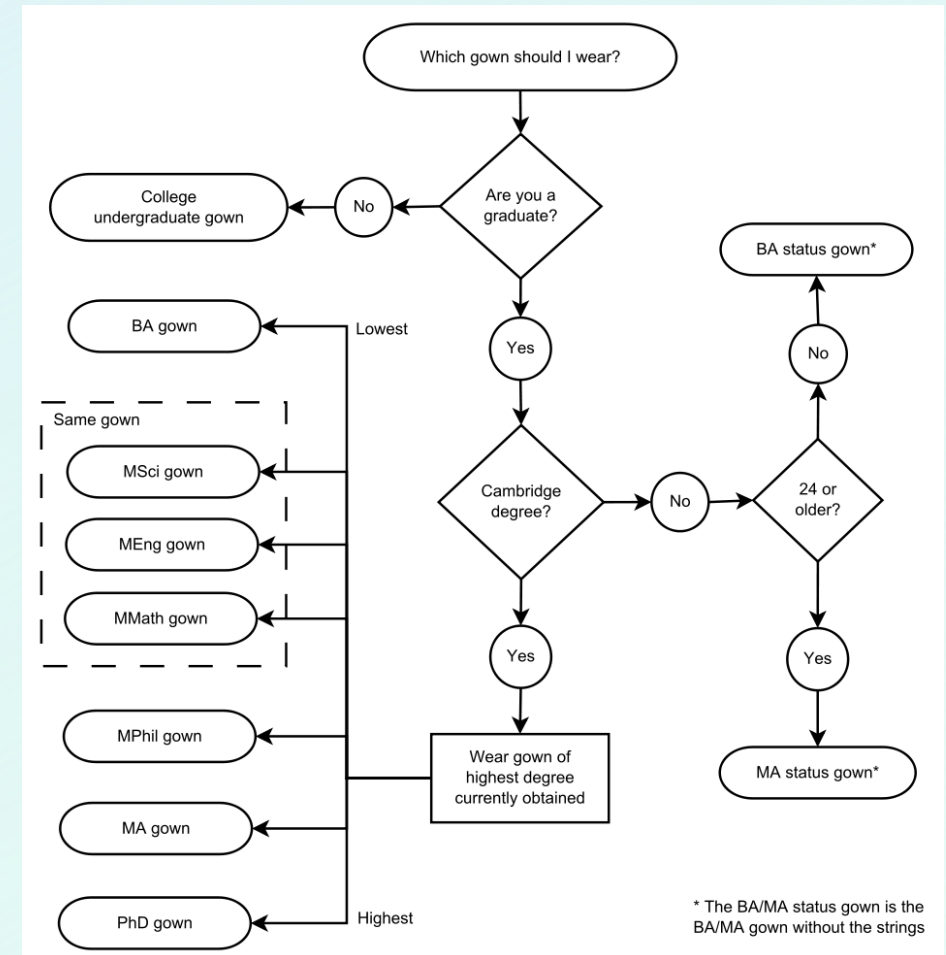


Knowledge-Based AI

Write programs that simulate how people solve the problem

Fundamental limitations:

- Will never get better than a person
- Requires deep domain knowledge
- We don't know how we do some things (e.g., riding a bicycle)



Data-Based AI = Machine Learning

Write **programs that learn** the task **from examples**

- ✓ No need to know how we do it as humans
- ✓ Performance should improve with more examples
- ✗ May need **many examples!**
- ✗ May not understand how the program works!

Machine Learning:

Study of algorithms that

- Improve their performance P
- At some task T
- With experience E

Well defined task: $\langle P, T, E \rangle$

The Machine Learning Process

Experience

- Examples of the form
(input, correct output)

Task

- Mapping from input to output

Performance

- "Loss function" that measures error w.r.t. desired outcome

Choices in ML Problem Formulation

Experience

- Examples of the form (input, correct output)

Task

- Mapping from input to output

Performance

- "Loss function" that measures error w.r.t. desired outcome

Loan Applications

- What historical examples do I have? What is a correct output?
- Predict probability of default? Loan decision? Credit score?
- Do I care more about minimizing False Positives? False negatives?

Machine Learning:

Study of algorithms that

- Improve their performance P Optimization, Evaluation
- At some task T Classification, regression, clustering
- With experience E Tabular, image, sequence

Well defined task: <P,T,E>

How will I rate “Chopin’s 5th Symphony”?

Song	Rating
Some nights	★★★★★
Skyfall	★
Comfortably numb	★★★
We are young	★★★★
...	...
...	...
Chopin’s 5 th	???

Classification: Three Elements

1. Data:

- x : data example with d attributes
- y : label of example (what you care about)

2. Classification **model**: a function $f_{(a,b,c,\dots)}$

- Maps from X to Y
- (a,b,c,\dots) are the **parameters**

3. **Loss** function:

- Penalizes the model's mistakes

Song	Rating
Some nights	★★★★★
Skyfall	★
Comfortably numb	★★★
We are young	★★★★★
...	...
...	...
Chopin's 5 th	???

Terminology Explanation

Song	Artist	Length	...	Rating
Some nights	Fun	4:23	...	★★★★★
Skyfall	Adele	4:00	...	★
Comf. Numb	Pink Floyd	6:13	...	★★★
We are young	Fun	3:50	...	★★★★★
...
...
Chopin's 5 th	Chopin	5:32	...	???


Data example = data instance

Attribute = feature = dimension

Label = target attribute

What is a “model”?

A **useful approximation** of the world



Typically, there are **many reasonable models** for the same data

Training a model = finding appropriate values for (a,b,c,...)

- An **optimization** problem
- “appropriate” = **minimizes the Loss (cost)** function
- We will focus on a common training algorithm later on

Classification Loss Function

- How unhappy are you with the answer that the model gave?
- $L_{0-1}(y, f(x)) = \begin{cases} 1 & \text{if: } y \neq f(x) \\ 0 & \text{otherwise} \end{cases}$
- **0-1 loss** function: intuitive but hard to optimize = train
- In practice, we use **approximations** of the 0-1 loss – getting warmer or getting colder



Why should this work at all?

The main theoretical basis of ML:

With a **sufficient amount of “similar” data**

+

an **expressive model class**:

Minimizing the loss function on the training data yields a highly **accurate model on unseen test data**, with high probability

1. **Data:** $S = \{(x_i, y_i)\}_{i=1, \dots, n}$

- x_i : data example with d attributes
- y_i : label of example (what you care about)

2. **Classification model:** a function $f_{(a,b,c,\dots)}$

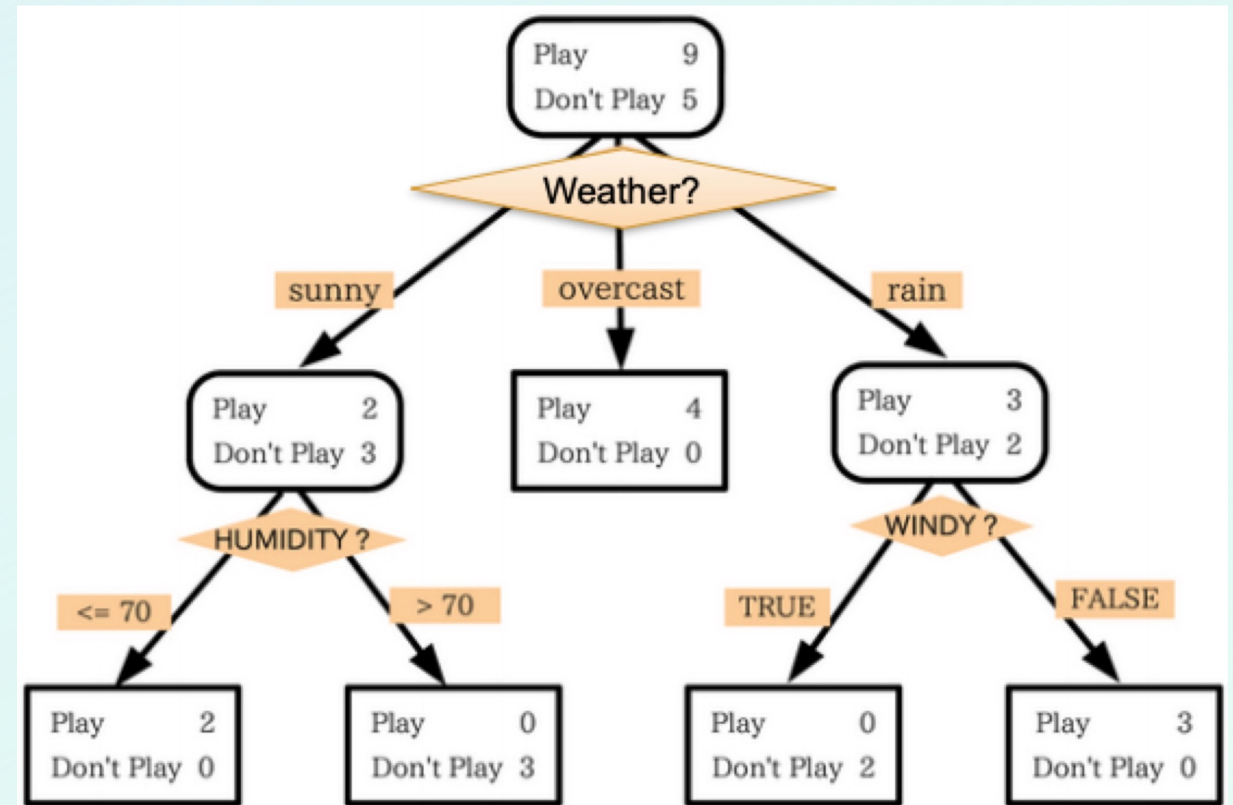
- Maps from X to Y
- (a,b,c,\dots) are the **parameters**

3. **Loss function:** $L(y, f(x))$

- Penalizes the model's mistakes

Decision Trees: To play **tennis** or not to?


- **Data:** attributes describing the weather; (sunny? humidity level, ...)
- **Target:** 1 if it's good to **Play**, 0 otherwise
- **Model:** $f_T(x)$
- **Model parameters:** T , the tree structure (and size)



Training (fitting) a Decision Tree


How to choose the attribute/value to split on at each level of the tree?

- Two classes (red circles / green crosses)
- Two attributes: X and Y
- 11 points in training data
- Idea: construct a decision tree such that the leaf nodes correctly predict the class for all the training examples




Training (fitting) a Decision Tree

These splits are great because the nodes are “**pure**”



These splits are bad because the nodes have a **mix** of samples




Training (=fitting) a Decision Tree

1. Find the **best attribute** to split on
2. Find the **best split** on the chosen attribute
3. Repeat 1 & 2 until **stopping criterion** is met

Common **stopping criteria**:

- Node contains very few data points
- Node is pure: most training data in node have same label



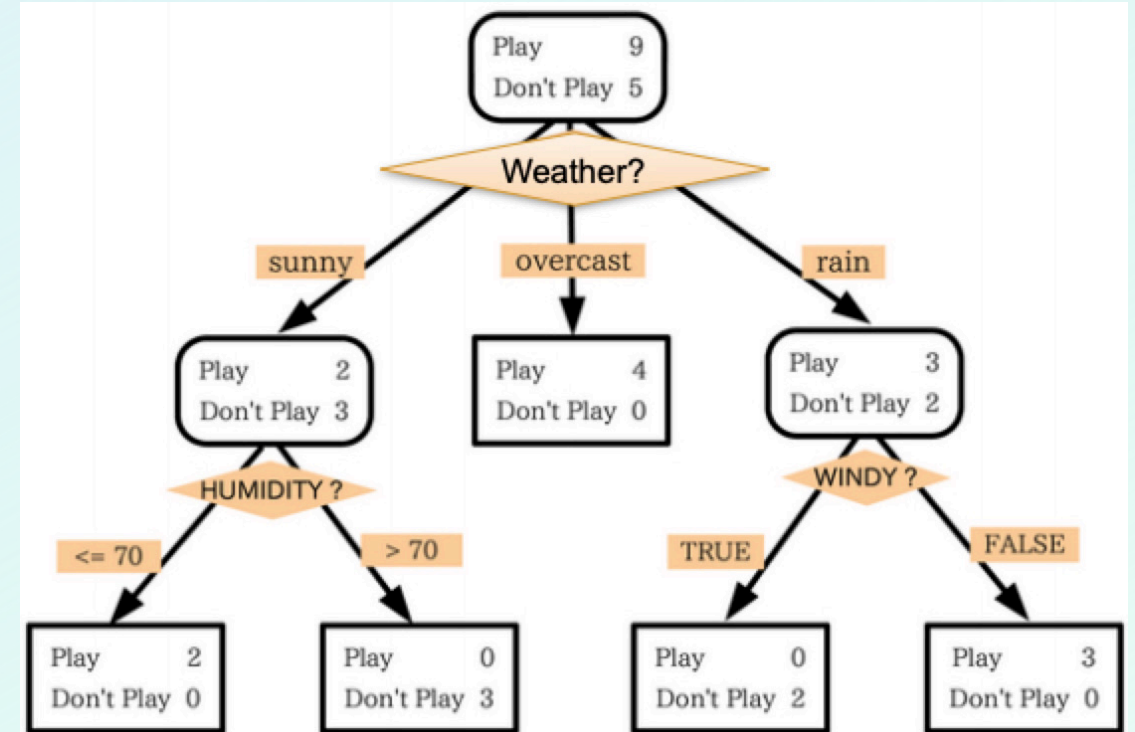
Final words on Decision Trees

Advantages

- Simple interpretation
- Fast predictions
- Handles mixed-type attributes

Caveats

- May be too simple for complex data
- Hard to figure out the right depth, stopping criterion, especially at the node level



Forecasting

Decision Trees predict **discrete** outcomes

- Given X features, what is y outcome?


What about something **continuous**?

- How is something likely to change over time?
- Time series data tracks a value over time
 - e.g. temperature in one location, company sales data
- Often we want to predict what will happen *next*, i.e. in sequence

Time Series Data: Three Main Components

I. Trend: the long-term trajectory of the data


Avocado production is steadily increasing year-on-year



Time Series Data: Three Main Components

2. Seasonality: how the data changes according to the calendar


This is highly domain specific – e-commerce platforms experience the most sales around Christmas, but homes are sold most frequently in the summer



Time Series Data: Three Main Components

3. Noise: The random element left over

Noise comes from short-term changes that are unpredictable. More noise makes it harder to anticipate change



Goals of Time Series Analysis

1. Describe the data: represent the entire dataset compactly
2. Interpretation: understand the factors that drive change in the data
3. Forecasting: Predict what will come next!